

Ola_Churn_Analysis

Python Internship Project-

Define Problem Statement and Perform Exploratory Data Analysis

Question 1-Definition of Problem:

Understand the challenge of driver attrition and its impact on Ola.

Answer-Many drivers leave Ola frequently or switch to Uber for better pay. This makes it hard for Ola to maintain a stable workforce. Finding and training new drivers costs a lot of money. It is cheaper to retain existing drivers than to keep hiring new ones. When too many drivers leave, it affects customer service, increases ride cancellations, and leads to longer wait times for passengers.

Impact on Ola: 1. Higher costs for hiring and training new drivers

2. Lower service quality due to driver shortages

3. Loss of experienced drivers, affecting customer satisfaction

4. Reduced profits because frequent hiring is expensive

[35]: Loading the data

```
import pandas as pd
df = pd.read_csv("ola_driver - ola_driver.csv")
```

[36]: *# Checking the data shape*
df.shape

[36]: (19104, 14)

Checking the data type of each column

[37]: *#checking the data type of each column*
df.dtypes

[37] : Unnamed: 0 int64
 MMM-YY object
 Driver_ID int64
 Age float64

```

Gender          float64
City            object
Education_Level int64
Income          int64
Dateofjoining  object
LastWorkingDate object
Joining Designation int64
Grade          int64
Total Business Value int64
Quarterly Rating int64
dtype: object

```

Conversion to text data types

```

[38] : # conversion to text data types
column_data_type_conversion = ["Driver_ID", "Gender", "Education_Level",
                                "Joining Designation", "Grade"]

df[column_data_type_conversion] = df[column_data_type_conversion].astype(str)

```

Detect missing values

```

[39] : #Detect missing values
df.isnull().sum()

```

```

[39] : Unnamed: 0          0
      MMM-YY           0
      Driver_ID        0
      Age             61
      Gender          0
      City            0
      Education_Level  0
      Income          0
      Dateofjoining    0
      LastWorkingDate 17488
      Joining Designation 0
      Grade           0
      Total Business Value 0
      Quarterly Rating  0
      dtype: int64

```

Perform statistical summary to understand data distribution.

```

[40] : #Perform statistical summary to understand data distribution.
df.describe()

```

```

[40] :
count    Unnamed: 0    Age    Income    Total Business Value \
mean    9551.500000    34.668435    65652.025126    5.716621e+05

```

std	5514.994107	6.257912	30914.515344	1.128312e+06
min	0.000000	21.000000	10747.000000	-6.000000e+06
25%	4775.750000	30.000000	42383.000000	0.000000e+00
50%	9551.500000	34.000000	60087.000000	2.500000e+05
75%	14327.250000	39.000000	83969.000000	6.997000e+05
max	19103.000000	58.000000	188418.000000	3.374772e+07

	Quarterly Rating
count	19104.000000
mean	2.008899
std	1.009832
min	1.000000
25%	1.000000
50%	2.000000
75%	3.000000
max	4.000000

Handle missing values using mean, median, or mode for numerical features.

```
[41] : #Handle missing values using mean, median, or mode for numerical features.
# filling null values of age column
df["Age"].fillna(df["Age"].mean(), inplace=True)
```

<ipython-input-41-a227f19ce7bc>:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["Age"].fillna(df["Age"].mean(), inplace=True)
```

Feature Engineering: Create a target variable indicating whether a driver has left the company based on LastWorkingDate.

```
[42] : #Feature Engineering: Create a target variable indicating whether a driver has
left the company based on LastWorkingDate.
df["Is_Driver_available"] = df["LastWorkingDate"].isnull().astype("int")
df.head()
```

/usr/local/lib/python3.11/dist-packages/google/colab/_dataframe_summarizer.py:88: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
cast_date_col = pd.to_datetime(column, errors="coerce")
```

```
[42]: Unnamed: 0    MMM-YY Driver_ID    Age Gender City    Education_Level    Income \
0          0    01/01/19          1  28.0    0.0  C23                2    57387
1          1    02/01/19          1  28.0    0.0  C23                2    57387
2          2    03/01/19          1  28.0    0.0  C23                2    57387
3          3    11/01/20          2  31.0    0.0   C7                2    67016
4          4    12/01/20          2  31.0    0.0   C7                2    67016
```

```
    Dateofjoining LastWorkingDate    Joining Designation Grade \
0      24/12/18                NaN                1        1
1      24/12/18                NaN                1        1
2      24/12/18      03/11/19                1        1
3      11/06/20                NaN                2        2
4      11/06/20                NaN                2        2
```

```
    Total Business Value    Quarterly Rating    Is_Driver_available
0          2381060                2                1
1         -665480                2                1
2              0                2                0
3              0                1                1
4              0                1                1
```

Calculate age of each driver based on Date Of Joining

```
[43] : #Calculate age of each driver based on Date Of Joining
```

```
df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"])

df["Years_after_joining"] = pd.Timestamp.today().year - df["Dateofjoining"].dt.
    year

df["Current Age"] = df["Age"] + df["Years_after_joining"]

df[["Driver_ID", "Age", "Dateofjoining", "Current Age"]].head()
```

<ipython-input-43-7b30dd0180f2>:3: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"])
```

```
[43] : Driver_ID    Age Dateofjoining    Current Age
0          1  28.0    2018-12-24        35.0
1          1  28.0    2018-12-24        35.0
2          1  28.0    2018-12-24        35.0
3          2  31.0    2020-11-06        36.0
```

4 2 31.0 2020-11-06 36.0

Determine if quarterly rating has increased (1 if yes, 0 if no).

```
[44] : #Determine if quarterly rating has increased (1 if yes, 0 if no).

df = df.sort_values(by=["Driver_ID", "MMM-YY"]) # firstly sorted the values in_
      □ ascending order.
df["Rating_Increased"] = df.groupby("Driver_ID")["Quarterly Rating"].diff().
      □ gt(0).astype(int)
df
```

/usr/local/lib/python3.11/dist-
 packages/google/colab/_dataframe_summarizer.py:88: UserWarning: Could not infer
 format, so each element will be parsed individually, falling back to `dateutil`.
 To ensure parsing is consistent and as-expected, please specify a format.
 cast_date_col = pd.to_datetime(column, errors="coerce")

```
[44]:      Unnamed: 0      MMM-YY Driver_ID   Age Gender City Education_Level  \
0              0  01/01/19           1  28.0    0.0  C23                2
1              1  02/01/19           1  28.0    0.0  C23                2
2              2  03/01/19           1  28.0    0.0  C23                2
6674          6674  01/01/20        1000  27.0    1.0   C3                2
6675          6675  02/01/20        1000  28.0    1.0   C3                2
...          ...      ...      ...      ...      ...      ...
6668          6668  08/01/20           999  23.0    1.0  C20                1
6669          6669  09/01/20           999  23.0    1.0  C20                1
6670          6670  10/01/20           999  23.0    1.0  C20                1
6671          6671  11/01/20           999  23.0    1.0  C20                1
6672          6672  12/01/20           999  23.0    1.0  C20                1
```

```
      Income Dateofjoining LastWorkingDate Joining  Designation Grade  \
0      57387    2018-12-24              NaN              1      1
1      57387    2018-12-24              NaN              1      1
2      57387    2018-12-24    03/11/19              1      1
6674    56016    2019-11-28              NaN              1      1
6675    56016    2019-11-28              NaN              1      1
...      ...      ...      ...      ...      ...
6668    36811    2020-05-05              NaN              2      2
6669    36811    2020-05-05              NaN              2      2
6670    36811    2020-05-05              NaN              2      2
6671    36811    2020-05-05              NaN              2      2
6672    36811    2020-05-05              NaN              2      2
```

```
      Total Business Value  Quarterly Rating  Is_Driver_available  \
0              2381060                2                1
1             -665480                2                1
2                0                2                0
```

6674	0	1	1
6675	0	1	1
...
6668	180170	2	1
6669	274440	2	1
6670	845830	2	1
6671	0	2	1
6672	247400	2	1

	Years_after_joining	Current Age	Rating_Increased
0	7	35.0	0
1	7	35.0	0
2	7	35.0	0
6674	6	33.0	0
6675	6	34.0	0
...
6668	5	28.0	0
6669	5	28.0	0
6670	5	28.0	0
6671	5	28.0	0
6672	5	28.0	0

[19104 rows x 18 columns]

Identify if monthly income has increased (1 if yes, 0 if no).

```
[45] : #Identify if monthly income has increased (1 if yes, 0 if no).
df["Month"] = pd.to_datetime(df["MMM-YY"]).dt.to_period("M") # Converting to_
           □ month format
df["Is_Income_Increased"] = df.groupby(["Driver_ID", "Month"])["Income"].diff().
           □ gt(0).astype(int)
df
```

<ipython-input-45-c2c982de9ffa>:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df["Month"] = pd.to_datetime(df["MMM-YY"]).dt.to_period("M") # Converting to
month format
```

/usr/local/lib/python3.11/dist-

packages/google/colab/_dataframe_summarizer.py:88: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`.

To ensure parsing is consistent and as-expected, please specify a format.

```
cast_date_col = pd.to_datetime(column, errors="coerce")
```

```
[45] : Unnamed: 0  MMM-YY  Driver_ID  Age  Gender  City  Education_Level  \
0      0  01/01/19      1  28.0    0.0  C23      2
1      1  02/01/19      1  28.0    0.0  C23      2
2      2  03/01/19      1  28.0    0.0  C23      2
```

6674	6674	01/01/20	1000	27.0	1.0	C3	2
6675	6675	02/01/20	1000	28.0	1.0	C3	2
...
6668	6668	08/01/20	999	23.0	1.0	C20	1
6669	6669	09/01/20	999	23.0	1.0	C20	1
6670	6670	10/01/20	999	23.0	1.0	C20	1
6671	6671	11/01/20	999	23.0	1.0	C20	1
6672	6672	12/01/20	999	23.0	1.0	C20	1

	Income	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
0	57387	2018-12-24	NaN			1	1
1	57387	2018-12-24	NaN			1	1
2	57387	2018-12-24	03/11/19			1	1
6674	56016	2019-11-28	NaN			1	1
6675	56016	2019-11-28	NaN			1	1
...
6668	36811	2020-05-05	NaN			2	2
6669	36811	2020-05-05	NaN			2	2
6670	36811	2020-05-05	NaN			2	2
6671	36811	2020-05-05	NaN			2	2
6672	36811	2020-05-05	NaN			2	2

	Total	Business Value	Quarterly	Rating	Is_Driver_available	\
0		2381060		2	1	
1		-665480		2	1	
2		0		2	0	
6674		0		1	1	
6675		0		1	1	
...
6668		180170		2	1	
6669		274440		2	1	
6670		845830		2	1	
6671		0		2	1	
6672		247400		2	1	

	Years_after_joining	Current Age	Rating_Increased	Month	\
0	7	35.0	0	2019-01	
1	7	35.0	0	2019-02	
2	7	35.0	0	2019-03	
6674	6	33.0	0	2020-01	
6675	6	34.0	0	2020-02	
...
6668	5	28.0	0	2020-08	
6669	5	28.0	0	2020-09	
6670	5	28.0	0	2020-10	
6671	5	28.0	0	2020-11	
6672	5	28.0	0	2020-12	

	Is_Income_Increased
0	0
1	0
2	0
6674	0
6675	0
...	...
6668	0
6669	0
6670	0
6671	0
6672	0

[19104 rows x 20 columns]


Check for class imbalance in the target variable.

```
[46] : #Check for class imbalance in the target variable.
Driver_available_vs_not_available=df["Is_Driver_available"].value_counts()
Driver_available_vs_not_available
```

```
[46] : Is_Driver_available
1      17488
0       1616
Name: count, dtype: int64
```

```
[47] : #Check for class imbalance in the target variable.
import matplotlib.pyplot as plt
plt.
    pie(Driver_available_vs_not_available,labels=Driver_available_vs_not_available.
    index)
```

```
[47] : ([<matplotlib.patches.Wedge at 0x7bac56045b50>,
    <matplotlib.patches.Wedge at 0x7bac53f42690>],
    [Text(-1.0613865244046894, 0.2888921006398996, '1'),
    Text(1.0613865496604358, -0.2888920078505379, '0')])
```

Insights from above graph-

Class Imbalance Observed-

1. There is a significant class imbalance in the Is_Driver_available variable.
2. 17,488 drivers (91.55%) are available, while 1,616 drivers (8.45%) are not available.

Potential Issues Due to Imbalance-

1. If used for a predictive model, the imbalance might cause biased results.
2. The model may favor the majority class (drivers available) and underperform in predicting driver churn.

What is the structure of the dataset (number of rows and columns)

```
[48] : #What is the structure of the dataset (number of rows and columns)  
df.shape
```

```
[48]: (19104, 20)
```

What are the data types of each column?

```
[49] : #What are the data types of each column?  
df.dtypes
```

```
[49] : Unnamed: 0          int64
      MMM-YY             object
      Driver_ID          object
      Age                float64
      Gender             object
      City               object
      Education_Level    object
      Income             int64
      Dateofjoining      datetime64[ns]
      LastWorkingDate    object
      Joining Designation object
      Grade              object
      Total Business Value int64
      Quarterly Rating   int64
      Is_Driver_available int64
      Years_after_joining int32
      Current Age         float64
      Rating_Increased    int64
      Month               period[M]
      Is_Income_Increased int64
      dtype: object
```

Are there any missing values in the dataset? If so, which columns are affected?

```
[50] : #Are there any missing values in the dataset? If so, which columns are affected?
      df.isnull().sum()
```

```
[50] : Unnamed: 0          0
      MMM-YY             0
      Driver_ID          0
      Age                0
      Gender             0
      City               0
      Education_Level    0
      Income             0
      Dateofjoining      0
      LastWorkingDate    17488
      Joining Designation 0
      Grade              0
      Total Business Value 0
      Quarterly Rating   0
      Is_Driver_available 0
      Years_after_joining 0
      Current Age         0
      Rating_Increased    0
      Month               0
      Is_Income_Increased 0
      dtype: int64
```

What are the basic statistics (mean, median, standard deviation) for numerical features like Age, Income, Total Business Value, and Quarterly Rating?

```
[51]: #What are the basic statistics (mean, median, standard deviation) for numerical_
      □ features like Age, Income, Total Business Value, and Quarterly Rating?
      column_data_type_conversion = ["Driver_ID", "Gender", "Education_Level",
      □ "Joining Designation", "Grade"]

      df[column_data_type_conversion] = df[column_data_type_conversion].astype(str)

      df.describe()
```

```
[51]:
```

	Unnamed: 0	Age	Income \
count	19104.000000	19104.000000	19104.000000
mean	9551.500000	34.668435	65652.025126
min	0.000000	21.000000	10747.000000
25%	4775.750000	30.000000	42383.000000
50%	9551.500000	34.000000	60087.000000
75%	14327.250000	39.000000	83969.000000
max	19103.000000	58.000000	188418.000000
std	5514.994107	6.247912	30914.515344

	Dateofjoining	Total Business Value	Quarterly Rating \
count	19104	1.910400e+04	19104.000000
mean	2018-04-28 20:52:54.874372096	5.716621e+05	2.008899
min	2013-04-01 00:00:00	-6.000000e+06	1.000000
25%	2016-11-29 12:00:00	0.000000e+00	1.000000
50%	2018-09-12 00:00:00	2.500000e+05	2.000000
75%	2019-11-05 00:00:00	6.997000e+05	3.000000
max	2020-12-28 00:00:00	3.374772e+07	4.000000
std	NaN	1.128312e+06	1.009832

	Is_Driver_available	Years_after_joining	Current Age \
count	19104.000000	19104.000000	19104.000000
mean	0.915410	7.223670	41.892105
min	0.000000	5.000000	26.000000
25%	1.000000	6.000000	37.000000
50%	1.000000	7.000000	41.000000
75%	1.000000	9.000000	47.000000
max	1.000000	12.000000	65.000000
std	0.278277	1.920872	7.038238

	Rating_Increased	Is_Income_Increased
count	19104.000000	19104.0
mean	0.146200	0.0
min	0.000000	0.0
25%	0.000000	0.0

50%	0.000000	0.0
75%	0.000000	0.0
max	1.000000	0.0
std	0.353316	0.0

How many unique drivers are there in the dataset?

```
[52] : #How many unique drivers are there in the dataset?
df["Driver_ID"].nunique()
```

[52]: 2381

How many drivers joined each month?

```
[53] : #How many drivers joined each month?
df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"]) # conversion of
    column to date format
df["LastWorkingDate"] = pd.to_datetime(df["LastWorkingDate"],errors="coerce")

df["Joining_Month"] = df["Dateofjoining"].dt.month # extraction of month from
    date
df["Leaving_month"] = df["LastWorkingDate"].dt.month

joins_per_month = df["Joining_Month"].value_counts().sort_index() # counting
    how many driver joined each month
exits_per_month = df["Leaving_month"].value_counts().sort_index()

joins_per_month
```

<ipython-input-53-04a2a9c86f3a>:3: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df["LastWorkingDate"] = pd.to_datetime(df["LastWorkingDate"],errors="coerce")
```

[53] : Joining_Month

```
1    1381
2     684
3     402
4    1014
5    2362
6    1973
7    2730
8    1886
9    1449
10   2095
11   1867
12   1261
```

Name: count, dtype: int64

How many drivers left each month?

```
[54] : df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"]) # conversion of_
      □column to date format
df["LastWorkingDate"] = pd.to_datetime(df["LastWorkingDate"],errors="coerce")

df["Joining_Month"] = df["Dateofjoining"].dt.month # extraction of month from_
      □date
df["Leaving_month"] = df["LastWorkingDate"].dt.month

joins_per_month = df["Joining_Month"].value_counts().sort_index() # counting_
      □how many driver left each month
left_per_month = df["Leaving_month"].value_counts().sort_index()

left_per_month
```

```
[54] : Leaving_month
1.0    152
2.0    155
3.0    133
4.0     91
5.0    161
6.0    138
7.0    189
8.0     57
9.0    145
10.0   132
11.0   142
12.0   121
Name: count, dtype: int64
```

Can we determine the average tenure of drivers in the dataset?

```
[55] : #Can we determine the average tenure of drivers in the dataset?

# Converting columns to datetime format
df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"], errors="coerce")
df["LastWorkingDate"] = pd.to_datetime(df["LastWorkingDate"], errors="coerce")

# Replacing missing values in LastWorkingDate with today's date.
df["LastWorkingDate"].fillna(pd.to_datetime("today"), inplace=True)

# Calculating total number of days drivers worked
df["Total_no_of_days"] = (df["LastWorkingDate"] - df["Dateofjoining"]).dt.days

# Calculating average days
average_tenure = df["Total_no_of_days"].mean()
```

```
average_tenure
```

<ipython-input-55-9d2bc88ea722>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["LastWorkingDate"].fillna(pd.to_datetime("today"), inplace=True)
```

```
[55]: np.float64(2362.453465242881)
```

Average tenure of drivers is 2361 days

How can we create a target variable to indicate whether a driver has left the company based on LastWorkingDate?

```
[56]: #How can we create a target variable to indicate whether a driver has left the_
      □company based on LastWorkingDate?
df["Is_Driver_available"]=df["LastWorkingDate"].isnull().astype("int")
df
```

```
[56]:
```

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	\
0	0	01/01/19	1	28.0	0.0	C23	2	
1	1	02/01/19	1	28.0	0.0	C23	2	
2	2	03/01/19	1	28.0	0.0	C23	2	
6674	6674	01/01/20	1000	27.0	1.0	C3	2	
6675	6675	02/01/20	1000	28.0	1.0	C3	2	
...	
6668	6668	08/01/20	999	23.0	1.0	C20	1	
6669	6669	09/01/20	999	23.0	1.0	C20	1	
6670	6670	10/01/20	999	23.0	1.0	C20	1	
6671	6671	11/01/20	999	23.0	1.0	C20	1	
6672	6672	12/01/20	999	23.0	1.0	C20	1	

	Income	Dateofjoining	LastWorkingDate	...	Quarterly Rating	\
0	57387	2018-12-24	2025-03-28	06:50:44.051884	...	2
1	57387	2018-12-24	2025-03-28	06:50:44.051884	...	2
2	57387	2018-12-24	2019-03-11	00:00:00.000000	...	2
6674	56016	2019-11-28	2025-03-28	06:50:44.051884	...	1
6675	56016	2019-11-28	2025-03-28	06:50:44.051884	...	1
...
6668	36811	2020-05-05	2025-03-28	06:50:44.051884	...	2

6669	36811	2020-05-05	2025-03-28	06:50:44.051884	...	2
6670	36811	2020-05-05	2025-03-28	06:50:44.051884	...	2
6671	36811	2020-05-05	2025-03-28	06:50:44.051884	...	2
6672	36811	2020-05-05	2025-03-28	06:50:44.051884	...	2

	Is_Driver_available	Years_after_joining	Current Age	Rating_Increased	\
0	0	7	35.0		0
1	0	7	35.0		0
2	0	7	35.0		0
6674	0	6	33.0		0
6675	0	6	34.0		0
...
6668	0	5	28.0		0
6669	0	5	28.0		0
6670	0	5	28.0		0
6671	0	5	28.0		0
6672	0	5	28.0		0

	Month	Is_Income_Increased	Joining_Month	Leaving_month	\
0	2019-01	0	12	NaN	
1	2019-02	0	12	NaN	
2	2019-03	0	12	3.0	
6674	2020-01	0	11	NaN	
6675	2020-02	0	11	NaN	
...
6668	2020-08	0	5	NaN	
6669	2020-09	0	5	NaN	
6670	2020-10	0	5	NaN	
6671	2020-11	0	5	NaN	
6672	2020-12	0	5	NaN	

	Total_no_of_days
0	2286
1	2286
2	77
6674	1947
6675	1947
...	...
6668	1788
6669	1788
6670	1788
6671	1788
6672	1788

[19104 rows x 23 columns]

What additional features can we extract from Dateofjoining, such as tenure or duration of employment?

```
[57] : #What additional features can we extract from Dateofjoining, such as tenure or_
      □duration of employment?
      # Converting Date of Joining column to datetime format
      df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"], errors="coerce")
      df["LastWorkingDate"] = pd.to_datetime(df["LastWorkingDate"], errors="coerce")

      # Fill missing LastWorkingDate with today's date for active drivers
      df["LastWorkingDate"].fillna(pd.to_datetime("today"), inplace=True)

      # Calculating "Total_no_of_days", "Total_no_of_months", "Total_no_of_years"
      □drivers worked
      df["Total_no_of_days"] = (df["LastWorkingDate"] - df["Dateofjoining"]).dt.days
      df["Total_no_of_months"] = df["Total_no_of_days"] // 30 # Approximate months
      df["Total_no_of_years"] = df["Total_no_of_days"] // 365 # Corrected years_
      □calculation

      df
```

<ipython-input-57-5c4d31a09ec6>:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["LastWorkingDate"].fillna(pd.to_datetime("today"), inplace=True)
```

[57]:

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	\
0	0	01/01/19	1	28.0	0.0	C23	2	
1	1	02/01/19	1	28.0	0.0	C23	2	
2	2	03/01/19	1	28.0	0.0	C23	2	
6674	6674	01/01/20	1000	27.0	1.0	C3	2	
6675	6675	02/01/20	1000	28.0	1.0	C3	2	
...	
6668	6668	08/01/20	999	23.0	1.0	C20	1	
6669	6669	09/01/20	999	23.0	1.0	C20	1	
6670	6670	10/01/20	999	23.0	1.0	C20	1	
6671	6671	11/01/20	999	23.0	1.0	C20	1	
6672	6672	12/01/20	999	23.0	1.0	C20	1	

	Income	Dateofjoining	LastWorkingDate	...	\
0	57387	2018-12-24	2025-03-28 06:50:44.051884	...	

1	57387	2018-12-24	2025-03-28	06:50:44.051884	...
2	57387	2018-12-24	2019-03-11	00:00:00.000000	...
6674	56016	2019-11-28	2025-03-28	06:50:44.051884	...
6675	56016	2019-11-28	2025-03-28	06:50:44.051884	...
...
6668	36811	2020-05-05	2025-03-28	06:50:44.051884	...
6669	36811	2020-05-05	2025-03-28	06:50:44.051884	...
6670	36811	2020-05-05	2025-03-28	06:50:44.051884	...
6671	36811	2020-05-05	2025-03-28	06:50:44.051884	...
6672	36811	2020-05-05	2025-03-28	06:50:44.051884	...

	Years_after_joining	Current Age	Rating_Increased	Month \
0	7	35.0	0	2019-01
1	7	35.0	0	2019-02
2	7	35.0	0	2019-03
6674	6	33.0	0	2020-01
6675	6	34.0	0	2020-02
...
6668	5	28.0	0	2020-08
6669	5	28.0	0	2020-09
6670	5	28.0	0	2020-10
6671	5	28.0	0	2020-11
6672	5	28.0	0	2020-12

	Is_Income_Increased	Joining_Month	Leaving_month	Total_no_of_days \
0	0	12	NaN	2286
1	0	12	NaN	2286
2	0	12	3.0	77
6674	0	11	NaN	1947
6675	0	11	NaN	1947
...
6668	0	5	NaN	1788
6669	0	5	NaN	1788
6670	0	5	NaN	1788
6671	0	5	NaN	1788
6672	0	5	NaN	1788


	Total_no_of_months	Total_no_of_years
0	76	6
1	76	6
2	2	0
6674	64	5
6675	64	5
...
6668	59	4
6669	59	4
6670	59	4

6671	59	4
6672	59	4

[19104 rows x 25 columns]


What are the distributions of Age, Income, and Total Business Value?

```
[58] : #What are the distributions of Age, Income, and Total Business Value?
#What are the distributions of Age?
plt.hist(df["Age"],bins=30,color="lightblue",edgecolor="black")
plt.show()
```



Insights from above graph- 1. The histogram shows that the majority of drivers fall within the 30 to 40-year-old range. 2. There are fewer drivers below 25 and above 45, suggesting that younger individuals might not prefer this job or may lack the required driving experience. 3. Similarly, drivers above 50 years are very few, indicating that older individuals may leave the workforce earlier due to physical strain or other factors.


```
[59] : # distribution of income
plt.hist(df["Income"],bins=30,color="lightblue",edgecolor="black")
plt.show()
```



Insights from above graph-

1. Most drivers earn between 50,000 - 75,000.
2. Very few earn above 1,25,000, showing an income gap.
3. The income distribution is right-skewed, meaning fewer high earners.

```
[60] : #What are the distributions of Total Business Value?  
plt.hist(df["Total Business Value"],bins=30,color="lightblue",edgecolor="black")  
plt.show()
```



Insights from above graph-

1. Most drivers have very low Total Business Value, clustered near zero.
2. A few drivers have very high values, creating a right-skewed distribution.
3. This suggests a few top performers contribute a significant share of the business.

How does Quarterly Rating vary across different drivers and time periods?

```
[61] : #How does Quarterly Rating vary across different drivers and time periods?
#conversion to datetime format
df["MMM-YY"] = pd.to_datetime(df["MMM-YY"])


#Extracting month from "MMM-YY" column
df["Month"] = df["MMM-YY"].dt.to_period("M")

# Calculating Average Quarter Rating for each month
rating_trend = df.groupby("Month")["Quarterly Rating"].mean()
# Line Graph
plt.figure(figsize=(12, 6))
plt.plot(rating_trend.index.astype(str), rating_trend.values, marker="o",
        linestyle="-", color="b", label="Avg Quarterly Rating")
```

```
plt.xlabel("Time (Year-Month)")
plt.ylabel("Average Quarterly Rating")
plt.title("Quarterly Rating Trend Over Time")
plt.legend()
plt.show()
```

<ipython-input-61-9df7aa66e31d>:3: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df["MMM-YY"] = pd.to_datetime(df["MMM-YY"])
```



How should missing values in LastWorkingDate be treated, considering it indicates whether a driver has left?

```
[62] : #How should missing values in LastWorkingDate be treated, considering it_
       □ indicates whether a driver has left?
       df["LastWorkingDate"].fillna("Present", inplace=True)
```

<ipython-input-62-3813c5cbbbdd>:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.


The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["LastWorkingDate"].fillna("Present", inplace=True)
```

Are there any trends or patterns in the monthly income or business value acquired?

```
[63] : #Are there any trends or patterns in the monthly income or business value_  
       acquired?  
#Converting "MMM-YY" column to datetime format  
df["Month"] = pd.to_datetime(df["MMM-YY"]) # Auto-detects format  
  
# Sort data in ascending order  
df = df.sort_values(by="Month")  
  
#calculating average income for each month  
monthly_income_trend = df.groupby("Month")["Income"].mean()  
  
# line chart  
plt.figure(figsize=(12, 5))  
plt.plot(monthly_income_trend, marker="o", linestyle="-", color="b",  
       label="Monthly Income")  
plt.xlabel("Month")  
plt.ylabel("Average Income")  
plt.title("Monthly Income Trend")  
plt.legend()  
plt.show()
```



Insights-

Strong Growth Phase (Early 2019 - Mid 2019): The monthly income showed a steady upward trend, indicating increasing demand and business expansion.

Decline & Fluctuations (Late 2019 - Early 2020): A noticeable dip in income occurred, possibly

due to seasonality, market shifts, or external factors like competition or economic slowdowns.


Recovery & Stability (Mid 2020 - 2021): The business recovered from the decline, showing resilience with stable income levels, despite some fluctuations.

Potential Pandemic Impact (2020): The fluctuations in 2020 could be linked to COVID-19 disruptions, affecting demand but eventually leading to a recovery phase.

Calculating Average of total business value for each month

```
[64] : #Calculating Average of total business value for each month
business_value_trend = df.groupby("Month")["Total Business Value"].mean()

# line plot
plt.figure(figsize=(12, 5))
plt.plot(business_value_trend, marker="o", linestyle="-", color="r",
        label="Total Business Value")
plt.xlabel("Month")
plt.ylabel("Total Business Value")
plt.title("Business Value Acquired Over Time")
plt.legend()
plt.show()
```



Insights-

Fluctuating Business Trends – Business value shows significant ups and downs, indicating variability in performance across months.


Notable Spikes in 2019 & 2020 – Sharp increases in business value suggest the impact of special events, promotions, or seasonal demand.

Sudden Declines Post-Growth – Rapid drops following peak months indicate that growth was not sustained over time.

Gradual Recovery Over Time – After declines, business value shows a slow yet steady improvement, though not reaching previous peak levels.

Is there a correlation between Age and Income?

```
[65] : #Is there a correlation between Age and Income?  
df.plot.scatter(x="Age",y="Income",title="Income vs Age")  
plt.show()
```



Insights-

Income Grows with Experience – Most people earn more between ages 30 to 45, likely as they gain experience and move up in their careers.

More Variation in Income at Older Ages – After 40, some people earn a lot while others earn much less, possibly due to different career paths, promotions, or industry shifts.

Lower Earnings in the Early Career Stage – People in their 20s generally earn less since they are just starting their careers and working in entry-level roles.

No Fixed Pattern Between Age and Income – Income doesn't just go up with age—other factors like job type, industry, and skills matter a lot too.

How do Education_Level affect Total Business Value?


```
[66] : #How do Education_Level and City affect Total Business Value?
education_impact_on_business_value = df.groupby("Education_Level")["Total_
    Business Value"].mean().sort_values(ascending=False)
education_impact_on_business_value
```

```
[66] : Education_Level
1    601287.867133
0    565410.657872
2    545364.175755
Name: Total Business Value, dtype: float64
```

Insights-

Higher Education Tends to Drive More Business Value-Individuals with Education Level 1 have the highest average total business value (601,288), followed by Level 0 (565,410), and Level 2 (545,364).

How do City affect Total Business Value?

```
[67] : #How do Education_Level and City affect Total Business Value?
city_impact_on_business_value = df.groupby("City")["Total Business Value"].
    mean().sort_values(ascending=False)
city_impact_on_business_value
```

```
[67]: City
C13    796263.075571
C29    736637.511111
C12    667282.310867
C26    661837.445339
C5     634855.975610
C16    632585.712271
C19    630978.151986
C14    607931.635802
C28    591406.778917
C24    584712.426710
C21    572684.776119
C27    572039.312977
C8     566328.539326
C6     566042.954545
C22    559749.431397
C4     556092.266436
C2     553365.084746
C15    553266.636005
C18    550106.250000
C10    540753.736559
C11    538549.145299
C1     531560.280650
C25    507575.119863
C7     484569.228243
C20    468535.605159
```

```
C9      467914.865385
C3      458003.940345
C17     429160.204545
C23     423986.561338
Name: Total Business Value, dtype: float64
```

Insights-

Different cities contribute significantly different average total business values, indicating that location plays a crucial role in business performance.

Are drivers with higher Quarterly Rating more likely to stay longer?

```
[70]: #Are drivers with higher Quarterly Rating more likely to stay longer?

# Converting dates to datetime format
df["Dateofjoining"] = pd.to_datetime(df["Dateofjoining"])
df["LastWorkingDate"] = pd.to_datetime(df["LastWorkingDate"])

# Filling missing LastWorkingDate with today's date
df["LastWorkingDate"].fillna(pd.to_datetime("today"), inplace=True)

# Calculating how many years the driver worked
df["Years Stayed"] = df["LastWorkingDate"].dt.year - df["Dateofjoining"].dt.year

#calculating average years for ach quarter rating
rating_trend = df.groupby("Quarterly Rating")["Years Stayed"].mean()

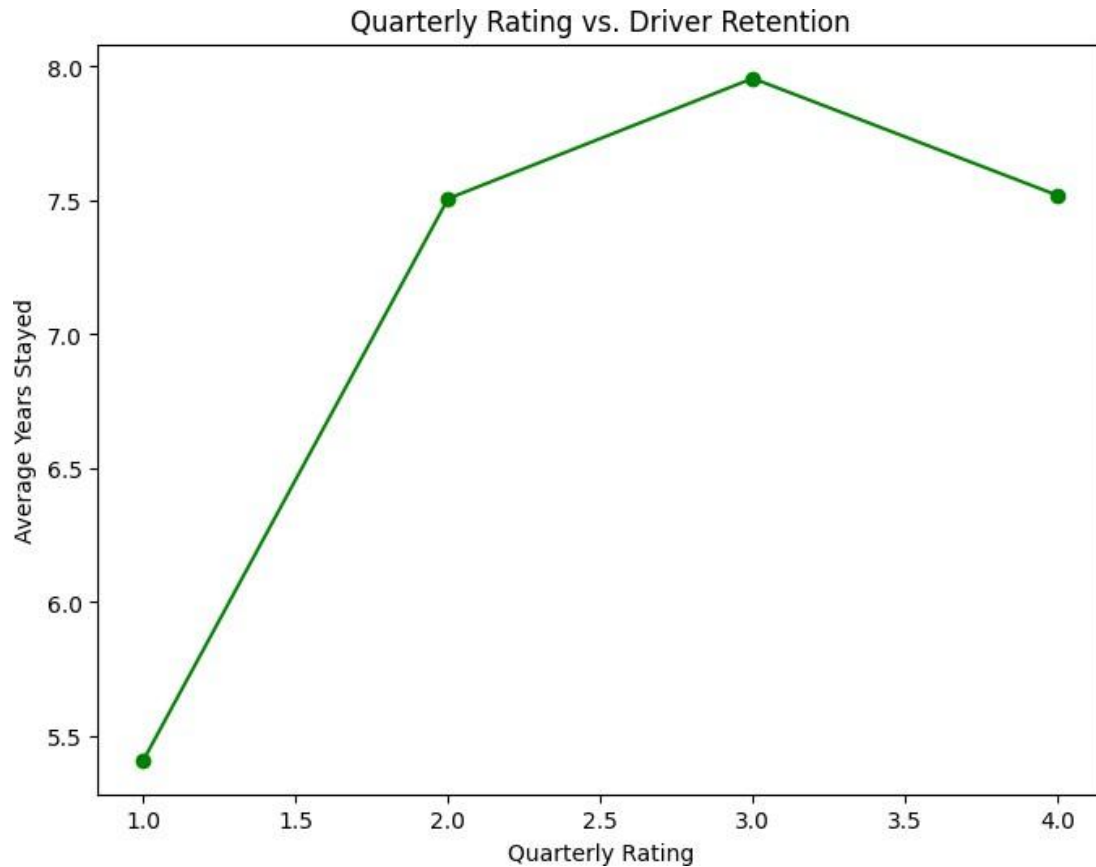
# line graph
plt.figure(figsize=(8,6))
plt.plot(rating_trend, marker="o", linestyle="--", color="g")
plt.xlabel("Quarterly Rating")
plt.ylabel("Average Years Stayed")
plt.title("Quarterly Rating vs. Driver Retention")
plt.show()
```

<ipython-input-70-ba28a9074b01>:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["LastWorkingDate"].fillna(pd.to_datetime("today"), inplace=True)
```



Insights-

The plot indicates that as the Quarterly Rating increases from 1 to 3, the Average Years Stayed also increases.

This suggests that drivers with higher ratings tend to stay longer, which makes sense as high-rated drivers may be more engaged, satisfied, or financially stable.

Interestingly, the trend declines at Rating 4, meaning drivers with the highest rating are leaving earlier than those with a rating of 3.

Highly-rated drivers might get better opportunities elsewhere.

Based on the analysis, what strategies can Ola implement to improve driver retention?

1. Ola should focus on increasing driver income, particularly for those earning in the 50,000 - 75,000 range, as many drivers leave for better pay
2. They should Implement performance-based incentives or bonuses to motivate drivers and increase their earning potential.

Are there specific demographic groups or performance metrics that require targeted interventions?

1. Drivers in the 30-40 age group represent the majority and should be a primary focus for retention strategies
2. Drivers earning in the 50,000 - 75,000 range may be more likely to seek alternative opportunities, so targeted interventions to increase their income are crucial
3. Drivers with consistently high Total Business Value are valuable assets and should be incentivized to stay with Ola.

Identify key factors influencing driver attrition.

1. Many drivers feel they're not making enough money due to reduced fares and high commission fees. If they can't cover their costs and make a decent living, they'll look for better opportunities elsewhere.
2. A single low rating or an unfair complaint can hurt a driver's standing on the platform, making it harder for them to get rides and earn money. Many feel this system is stacked against them.
3. Other ride-hailing services often offer better bonuses or lower commission cuts. Drivers will naturally move to where they feel more valued and better compensated.

Recommend strategies to improve driver retention.

1. Ensure that drivers are paid fairly, with bonuses for peak-hour work and reduced commission fees so they can actually take home a decent income.
2. Introduce rest periods and encourage work-life balance.
3. The company could focus on understanding why highly-rated drivers leave and implement targeted retention strategies (e.g., better incentives, promotions, or bonuses).
4. Reward drivers based on tenure with perks like lower fees, special discounts, or exclusive offers.
5. Develop retention programs specifically for drivers in the 30-40 age group, as they form the majority of the driver base
6. Allow drivers to choose shifts that better fit their personal lives. Avoid forcing excessive hours.